"Pull" - Mechanism	Advantages	Drawbacks	When to use
Data Pipelines	+ Out of the box connectors + One-To-One copy at scale + Orchestrating the whole process + Generic approach to private data sources	- Limited number of connectors - (almost) no transformation - Batch process	Large Scale ingestion If Connector is available
Data Flows Gen2	+ A lot of out of the box connectors + Transformation on the fly + Power Query Experience + Generic approach to private data sources	- (not yet) very performant - Batch process	Low-Code, small/medium scale ingestion and/or transformation
(End-to-End) Mirroring	+ Out-of-the-box near real-time replication + very easy set-up (for some sources)	- no transformation - very limited sources out of the box - currently no support for private data sources	Whenever possible and applicable (& not Data Lake/Cloud storage)
One Lake Shortcuts	+ No Replication of data + always up-to-date data	- only working for specific data stores - Potentially egress costs	Whenever possible and applicable (Data Lake/Cloud storage)
Spark (Notebooks, Spark Job Definition)	+ Large scale ingestion and/or transformation + Full flexibility of connectors and transformations via libraries	- Only limited private data sources possible - No Fabric-managed connectors	Code first, large scale batch or streaming ingestion and transformation If you don´t care about managed connectors
Eventstream	+ managed stream ingestion + out of the box connectors to streaming sources	- Limited to streaming	Ingest from streaming data sources via a managed service

"Push"- Mechanism	How?	When to use? (Generally when preferred Pull-Mechanisms are not applicable or external Tool is preferred)
Write to One Lake	One Lake supports the same API as Azure Data Lake Gen 2, i.e. you can connect via BLOB API: https://asdfsadfasf.onelake. ABFSS API: abfss://asdfasdfsadf Azure Storage Explorer and every tool using those APIs, e.g. Azure Databricks or Azure Data Factory	This is the least (Fabric-)managed way but also the most flexible one in the sense of being open to use any Data Ingest Tooling
Open Mirroring	Fabric offers the framework to do mirroring for any kind of data source as long as you land the changed data in the landing zone of Mirrored Database item.	Whenever applicable (if solution is already there or can be built with reasonable effort)
SQL Endpoint Ingest	The SQL Endpoints of SQL Databases and Warehouses in Fabric can not only be used for read but also write operations and can be a target of ETL-Tools targeting a MS SQL Endpoint.	SQL Database use cases (transactional or translytical)
Eventhouse Ingest	Kusto Databases which are offered as Azure Data Explorer on Azure expose APIs to ingest data. There are also libraries e.g. for PySpark to ingest data to a Kusto Database	Real-Time-Intelligence Use Cases when not using Eventstreams
Eventstream Endpoint	You can configure a Eventstream to have an Azure Eventhub endpoint and send events directl to this endpoint using the Eventhub SDKs.	General Real-Time-Intelligence Use Cases